

Throughput-Optimal Algorithms for Generalized Network-Flow Problems

Abhishek Sinha

Laboratory for Information and Decision Systems

Talk at: [TIFR](#)

December 29, 2016



Outline

- 1 Introduction
- 2 Optimal Broadcasting in a DAG
- 3 Algorithm for Generalized Flow
- 4 Simulation Results
- 5 Conclusion

Outline


- 1 Introduction
- 2 Optimal Broadcasting in a DAG
- 3 Algorithm for Generalized Flow
- 4 Simulation Results
- 5 Conclusion

Introduction

- We consider the **Generalized Flow Problem**, where the packets are to be routed in a wireless network from their source to the destination(s) at the **maximum rate** possible.
- Packets may belong to different sessions including **unicast**, **broadcast**, **multicast** or **anycast**.
- **A Fundamental** problem with wide ranging applications: Internet routing, in-network function computations, live multi-media streaming, military communications
- Contributions of my thesis:
 - **Broadcast**: A specialized **dynamic** algorithm that solves the optimal broadcasting problem
 - **Generalized Flow**: A single general algorithmic paradigm that solves *all flow problems*.

Network Model


- **Multi-hop** wireless network given by the graph $\mathcal{G}(V, E)$.
- **Time-slotted** system
- Stochastic arrivals- i.i.d. process with arrival vector λ
 - λ is, in general, **unknown** in advance



The network $\mathcal{G}(V, E)$

Wireless Interference Model


- Due to Wireless interference, only a **subset of links** can be activated at a slot. The set of all feasible link activations is given by \mathcal{M} .
 - **Example:** (1) For **primary interference** constraints: \mathcal{M} is the set of all **Matchings**
 - (2) For **wired** network : \mathcal{M} is the set of all subsets of links (no interference)




A wireless network

Wireless Interference Model

- Due to Wireless interference, only a **subset of links** can be activated at a slot. The set of all feasible link activations is given by \mathcal{M} .
 - Example:** (1) For **primary interference** constraints: \mathcal{M} is the set of all **Matchings**
 - (2) For **wired** network : \mathcal{M} is the set of all subsets of links (no interference)




A wireless network



Activation s_1

Wireless Interference Model

- Due to Wireless interference, only a **subset of links** can be activated at a slot. The set of all feasible link activations is given by \mathcal{M} .
 - Example:** (1) For **primary interference** constraints: \mathcal{M} is the set of all **Matchings**
 - (2) For **wired** network : \mathcal{M} is the set of all subsets of links (no interference)




A wireless network




Activation s_1 Activation s_2

Wireless Interference Model

- Due to Wireless interference, only a **subset of links** can be activated at a slot. The set of all feasible link activations is given by \mathcal{M} .
 - Example:** (1) For **primary interference** constraints: \mathcal{M} is the set of all **Matchings**
 - (2) For **wired** network : \mathcal{M} is the set of all subsets of links (no interference)



A wireless network

Activation s_1 Activation s_2 Activation s_3

A wireless network and its feasible link activations under **primary** interference constraints

- The links may also be time-varying $\{\text{ON}, \text{OFF}\}$ according to a stationary ergodic process.

Classification of Network Flow Problems

Traffic Classes: Traffic Class c has arrival rate λ_c , source node S^c , destination node(s) D^c , where

- **Unicast:** Single source, single destination

$$\mathcal{S} = \{s\}, \mathcal{D} = \{d\}$$

- **Multicast:** Single source, multiple destinations

$$\mathcal{S} = \{s\}, \mathcal{D} \subset V \setminus \{s\}$$

- **Broadcast:** Single source, all destinations

$$\mathcal{S} = \{s\}, \mathcal{D} = V \setminus \{s\}$$

- **Anycast:** Single source, choice of one among multiple alternative destinations

$$\mathcal{S} = \{s\}, \mathcal{D} = v_1 \oplus v_2 \oplus \dots \oplus v_k$$

- All of the above problems with multiple sources

Problem Statement

Design a **routing and scheduling** policy that supports all arrival rates λ within the network stability region.

- Including an arbitrary mix of unicast, multicast, broadcast traffic

Formally, let $R_{\pi}^{(c)}(T)$ denote the number of packets received by all destinations of class c upto time T under the policy π . Our objective is to find a policy π such that the following holds:

$$\liminf_{T \nearrow \infty} \frac{R_{\pi}^{(c)}(T)}{T} = \lambda^c, \quad \forall c \in \mathcal{C} \quad \text{w.p. } 1$$

For **all** arrival rates λ in the interior of the stability region Λ of the network.

Problem Statement

Design a **routing and scheduling** policy that supports all arrival rates λ within the network stability region.

- Including an arbitrary mix of unicast, multicast, broadcast traffic

Formally, let $R_{\pi}^{(c)}(T)$ denote the number of packets received by all destinations of class c upto time T under the policy π . Our objective is to find a policy π such that the following holds:

$$\liminf_{T \nearrow \infty} \frac{R_{\pi}^{(c)}(T)}{T} = \lambda^c, \quad \forall c \in \mathcal{C} \quad \text{w.p. } 1$$

For **all** arrival rates λ in the interior of the stability region Λ of the network.

We will design **two distinct optimal policies**.

- First, we consider the **Broadcast** problem on a Wireless Directed Acyclic Graph (**DAG**)
- Finally, we will consider the **general problem** stated above

Outline

- 1 Introduction
- 2 Optimal Broadcasting in a DAG
- 3 Algorithm for Generalized Flow
- 4 Simulation Results
- 5 Conclusion

The Broadcast Capacity λ^*

Observation : From source r to each node $t \neq r$,

$$\lambda^* \leq \text{MAX-FLOW}(r \rightarrow t)$$

The Broadcast Capacity λ^*

Observation : From source r to each node $t \neq r$,

$$\lambda^* \leq \text{MAX-FLOW}(r \rightarrow t)$$

Thus,

$$\lambda^* \leq \min_{t \in V \setminus \{r\}} \text{MAX-FLOW}(r \rightarrow t)$$

The Broadcast Capacity λ^*

Observation : From source r to each node $t \neq r$,

$$\lambda^* \leq \text{MAX-FLOW}(r \rightarrow t)$$

Thus,

$$\lambda^* \leq \min_{t \in V \setminus \{r\}} \text{MAX-FLOW}(r \rightarrow t)$$

Edmond's Tree-Packing Theorem [1965]

- The above inequality is satisfied with equality.
- There exist λ^* **edge-disjoint spanning trees** to achieve the broadcast capacity.

The Broadcast Capacity λ^*

Observation : From source r to each node $t \neq r$,


$$\lambda^* \leq \text{MAX-FLOW}(r \rightarrow t)$$

Thus,

$$\lambda^* \leq \min_{t \in V \setminus \{r\}} \text{MAX-FLOW}(r \rightarrow t)$$

Edmond's Tree-Packing Theorem [1965]

- The above inequality is satisfied with equality.
- There exist λ^* **edge-disjoint spanning trees** to achieve the broadcast capacity.



The Broadcast Capacity λ^*

Observation : From source r to each node $t \neq r$,


$$\lambda^* \leq \text{MAX-FLOW}(r \rightarrow t)$$

Thus,

$$\lambda^* \leq \min_{t \in V \setminus \{r\}} \text{MAX-FLOW}(r \rightarrow t)$$

Edmond's Tree-Packing Theorem [1965]

- The above inequality is satisfied with equality.
- There exist λ^* **edge-disjoint spanning trees** to achieve the broadcast capacity.



The Broadcast Capacity λ^*

Observation : From source r to each node $t \neq r$,


$$\lambda^* \leq \text{MAX-FLOW}(r \rightarrow t)$$

Thus,

$$\lambda^* \leq \min_{t \in V \setminus \{r\}} \text{MAX-FLOW}(r \rightarrow t)$$


Edmond's Tree-Packing Theorem [1965]

- The above inequality is satisfied with equality.
- There exist λ^* **edge-disjoint spanning trees** to achieve the broadcast capacity.




Prior Works

- Pre-computes the set of all spanning trees offline in wireline networks [Sarkar and Tassiulas, Trans. on IT 2002].



Prior Works

- Pre-computes the set of all spanning trees offline in wireline networks [Sarkar and Tassiulas, Trans. on IT 2002].



- Impractical for large and time-varying networks. ✗
- Wireless case is studied with a fixed activation schedule only [Towsley and Twigg, 2008].

Feasible Policy Space Π

A feasible broadcast policy $\pi \in \Pi$ executes following two actions at every slot t :

- **Link Activation $\pi(\mathcal{A})$:** Activate a subset of links (e.g., a **matching**) subject to the underlying interference constraints.
- **Packet Scheduling $\pi(\mathcal{S})$:** Transmit packets over the set of activated links.

Feasible Policy Space Π

A feasible broadcast policy $\pi \in \Pi$ executes following two actions at every slot t :


- **Link Activation $\pi(\mathcal{A})$:** Activate a subset of links (e.g., a **matching**) subject to the underlying interference constraints.
- **Packet Scheduling $\pi(\mathcal{S})$:** Transmit packets over the set of activated links.

Feasible Policy Space Π

A feasible broadcast policy $\pi \in \Pi$ executes following two actions at every slot t :

- **Link Activation $\pi(\mathcal{A})$:** Activate a subset of links (e.g., a **matching**) subject to the underlying interference constraints.
- **Packet Scheduling $\pi(\mathcal{S})$:** Transmit packets over the set of activated links.

- An arbitrary $\pi(\mathcal{S})$ is **hard** to describe : state-space¹ grows **exponentially**!




1. # of packets present at each subset of nodes

Feasible Policy Space Π

A feasible broadcast policy $\pi \in \Pi$ executes following two actions at every slot t :

- **Link Activation $\pi(\mathcal{A})$:** Activate a subset of links (e.g., a **matching**) subject to the underlying interference constraints.
- **Packet Scheduling $\pi(S)$:** Transmit packets over the set of activated links.

- An arbitrary $\pi(S)$ is **hard** to describe : state-space¹ grows **exponentially**!




1. # of packets present at each subset of nodes

Policy-space $\Pi^{\text{in-order}} \subset \Pi$

- To simplify $\pi(S)$, we consider the sub-space $\Pi^{\text{in-order}} \subset \Pi$, in which all packets are delivered to every node *in-order*.


Policy-space $\Pi^{\text{in-order}} \subset \Pi$

- To simplify $\pi(S)$, we consider the sub-space $\Pi^{\text{in-order}} \subset \Pi$, in which all packets are delivered to every node *in-order*.




Policy-space $\Pi^{\text{in-order}} \subset \Pi$

- To simplify $\pi(\mathcal{S})$, we consider the sub-space $\Pi^{\text{in-order}} \subset \Pi$ in which all packets are delivered to every node *in-order*.



Policy-space $\Pi^{\text{in-order}} \subset \Pi$

- To simplify $\pi(\mathcal{S})$, we consider the sub-space $\Pi^{\text{in-order}} \subset \Pi$ in which all packets are delivered to every node *in-order*.



Thus, the network state can be compactly represented by the vector

$$\mathbf{S}(t) = \{R_1(t), R_2(t), \dots, R_n(t)\}$$

Where $R_i(t)$ is the **total number** of packets received by node i up to time t .

Policy sub-space Π^*

Question: Does the reduced policy-space $\Pi^{\text{in-order}}$ contain an optimal policy?

Policy sub-space Π^*

Question: Does the reduced policy-space $\Pi^{\text{in-order}}$ contain an optimal policy?

We show that $\Pi^{\text{in-order}}$ **does contain** a **optimal** policy for **DAGs**.

To prove this, we consider a restricted subspace $\Pi^* \subset \Pi^{\text{in-order}}$ defined as follows :

Policy sub-space Π^*

Question: Does the reduced policy-space $\Pi^{\text{in-order}}$ contain an optimal policy?

We show that $\Pi^{\text{in-order}}$ **does contain** a **optimal** policy for **DAGs**.

To prove this, we consider a restricted subspace $\Pi^* \subset \Pi^{\text{in-order}}$ defined as follows :

$$\Pi^* \subset \Pi^{\text{in-order}} \subset \Pi$$

For all $\pi \in \Pi^*$, a packet p is eligible for transmission to node n iff **all** in-neighbors of node n contain the packet p .

Policy sub-space Π^*


Question: Does the reduced policy-space $\Pi^{\text{in-order}}$ contain an optimal policy?

We show that $\Pi^{\text{in-order}}$ **does contain** a **optimal** policy for **DAGs**.

To prove this, we consider a restricted subspace $\Pi^* \subset \Pi^{\text{in-order}}$ defined as follows :

$$\Pi^* \subset \Pi^{\text{in-order}} \subset \Pi$$

For all $\pi \in \Pi^*$, a packet p is eligible for transmission to node n iff **all** in-neighbors of node n contain the packet p .




Packet 1 → node d ✓

Packet 2 → node d ✗


In Summary: Policy hierarchies

All policies are **un-restricted** for link-activations ($\pi(\mathcal{A})$)



In Summary: Policy hierarchies


All policies are **un-restricted** for link-activations ($\pi(\mathcal{A})$)



Π : all policies that perform **arbitrary** packet-forwarding

In Summary: Policy hierarchies

All policies are **un-restricted** for link-activations ($\pi(\mathcal{A})$)




Π : all policies that perform **arbitrary** packet-forwarding

$\Pi^{\text{in-order}}$: policies that enforce **in-order** packet-forwarding

In Summary: Policy hierarchies

All policies are **un-restricted** for link-activations ($\pi(\mathcal{A})$)




Π : all policies that perform **arbitrary** packet-forwarding

$\Pi^{\text{in-order}}$: policies that enforce **in-order** packet-forwarding

Π^* : policies that allow reception only if **all in-neighbors** have received the specific packet

In Summary: Policy hierarchies

All policies are **un-restricted** for link-activations ($\pi(\mathcal{A})$)



Π : all policies that perform **arbitrary** packet-forwarding

$\Pi^{in-order}$: policies that enforce **in-order** packet-forwarding

Π^* : policies that allow reception only if **all in-neighbors** have received the specific packet

Punchline : There exists $\pi^* \in \Pi^*$, which is optimal for DAG!

System-dynamics under Π^*

State Variables:

For each node $j \in V \setminus \{r\}$ define

$$X_j(t) = \min_{i:(i,j) \in E} \left(R_i(t) - R_j(t) \right) \quad (\text{Relative deficiency})$$

$$i_t^*(j) = \arg \min_{i:(i,j) \in E} \left(R_i(t) - R_j(t) \right)$$

System-dynamics under Π^*

State Variables:

For each node $j \in V \setminus \{r\}$ define

$$X_j(t) = \min_{i:(i,j) \in E} \left(R_i(t) - R_j(t) \right) \quad (\text{Relative deficiency})$$

$$i_t^*(j) = \arg \min_{i:(i,j) \in E} \left(R_i(t) - R_j(t) \right)$$

Theorem 1: Queuing Dynamics

For any policy $\pi \in \Pi^*$ the dynamics of the state variables $\{\mathbf{X}(t)\}_{t \geq 0}$ satisfies a **Lindley recursion**.

System-dynamics under Π^*

State Variables:

For each node $j \in V \setminus \{r\}$ define

$$X_j(t) = \min_{i:(i,j) \in E} \left(R_i(t) - R_j(t) \right) \quad (\text{Relative deficiency})$$

$$i_t^*(j) = \arg \min_{i:(i,j) \in E} \left(R_i(t) - R_j(t) \right)$$

Theorem 1: Queuing Dynamics

For any policy $\pi \in \Pi^*$ the dynamics of the state variables $\{\mathbf{X}(t)\}_{t \geq 0}$ satisfies a **Lindley recursion**.

Theorem 2: Stability \implies Efficiency

Under Π^* , any algorithm stabilizing $\mathbf{X}(t)$ is a broadcast algorithm in a DAG.

Intuition : The state-vector $\mathbf{X}(t)$ mathematically corresponds to “queue-sizes” in the traditional queuing network.

Max-Weight Policy for Stabilizing $\mathbf{X}(t)$

Consider a policy $\pi^* \in \Pi^*$, for which $\pi^*(\mathcal{A})$ is obtained by minimizing the drift of the Lyapunov function $L(\mathbf{X}(t)) = \sum_{j \in V \setminus \{r\}} X_j^2(t)$

- To each edge $(i, j) \in E$, assign a weight $W_{ij}(t)$, where

$$W_{ij}(t) = \left(X_j(t) - \sum_{k: j = i_t^*(k)} X_k(t) \right) \mathbb{1}(\sigma_{ij}(t) = 1)$$

Max-Weight Policy for Stabilizing $\mathbf{X}(t)$

Consider a policy $\pi^* \in \Pi^*$, for which $\pi^*(\mathcal{A})$ is obtained by minimizing the drift of the Lyapunov function $L(\mathbf{X}(t)) = \sum_{j \in V \setminus \{r\}} X_j^2(t)$

- To each edge $(i, j) \in E$, assign a weight $W_{ij}(t)$, where

$$W_{ij}(t) = \left(X_j(t) - \sum_{k: j = i_t^*(k)} X_k(t) \right) \mathbb{1}(\sigma_{ij}(t) = 1)$$

- Choose a *max-weight* activation with weights $\mathbf{W}(t)$

$$\mu^{\pi^*}(t) \in \arg \max_{\mu \in \mathcal{C} \odot S} \mu \cdot \mathbf{W}(t)$$

Max-Weight Policy for Stabilizing $\mathbf{X}(t)$

Consider a policy $\pi^* \in \Pi^*$, for which $\pi^*(\mathcal{A})$ is obtained by minimizing the drift of the Lyapunov function $L(\mathbf{X}(t)) = \sum_{j \in V \setminus \{r\}} X_j^2(t)$

- To each edge $(i, j) \in E$, assign a weight $W_{ij}(t)$, where

$$W_{ij}(t) = \left(X_j(t) - \sum_{k: j=i_t^*(k)} X_k(t) \right) \mathbb{1}(\sigma_{ij}(t) = 1)$$

- Choose a *max-weight* activation with weights $\mathbf{W}(t)$

$$\mu^{\pi^*}(t) \in \arg \max_{\mu \in \mathcal{C} \odot \mathcal{S}} \mu \cdot \mathbf{W}(t)$$

Theorem 3: Throughput-Optimality of π^*

The policy π^* is an optimal broadcast policy for a DAG, i.e. for $\lambda < \lambda^*$

$$\liminf_{t \rightarrow \infty} \frac{R_i^{\pi^*}(t)}{t} = \lambda, \quad \forall i \in V \text{ w.p. 1}$$

Characterization of Broadcast-Capacity of Wireless DAGs

The policy yields constructive proof of the following:

Theorem 4: Capacity characterization

$$\lambda_{\text{DAG}}^* = \max_{\beta_\sigma \in \text{conv}(\mathcal{M}_\sigma)} \min_j \mathbf{u}_j \cdot \left(\sum_{\sigma} p(\sigma) \beta_\sigma \right)$$

where \mathbf{u}_j is the cut-vector separating the node j from the rest of the network.

Characterization of Broadcast-Capacity of Wireless DAGs

The policy yields constructive proof of the following:

Theorem 4: Capacity characterization

$$\lambda_{\text{DAG}}^* = \max_{\beta_\sigma \in \text{conv}(\mathcal{M}_\sigma)} \min_j \mathbf{u}_j \cdot \left(\sum_{\sigma} p(\sigma) \beta_\sigma \right)$$

where \mathbf{u}_j is the cut-vector separating the node j from the rest of the network.

This yields:

Theorem 5: Efficient Computation

λ_{DAG}^* may be computed in poly-time for **primary** interference constraint and **polynomially many** network-configurations.

The proof uses an LP formulation of the above and the well-known **poly-time separating oracle** for the **matching polytope**.

Summary of the results so far ...

- We focussed exclusively on the Broadcast Problem and designed a Throughput-optimal policy for Wireless DAGs
- The algorithm used the idea of *in-order* delivery of packets in a cleverly chosen *restricted* policy space.
- The Wireless *Broadcast capacity* of DAGs is *explicitly characterized*.
- We have also designed different throughput-optimal broadcast policy for arbitrary networks, using the new concept of *Backpressure on sets*. (not discussed in this talk)
- In the rest of the talk, we will consider the *generalized network flow problem*.

Outline

- 1 Introduction
- 2 Optimal Broadcasting in a DAG
- 3 Algorithm for Generalized Flow**
- 4 Simulation Results
- 5 Conclusion

What's new? (Preview)

The new policy is named **Universal Max-Weight** Algorithm (UMW). It is different from the classical algorithms (such as, *Backpressure*) in the following aspects:

What's new? (Preview)

The new policy is named **Universal Max-Weight** Algorithm (UMW). It is different from the classical algorithms (such as, *Backpressure*) in the following aspects:

- Instead of making routing decisions for each packet **hop-by-hop**, UMW **dynamically chooses** routes of each packet **right** at the **source**.

What's new? (Preview)

The new policy is named **Universal Max-Weight** Algorithm (UMW). It is different from the classical algorithms (such as, *Backpressure*) in the following aspects:

- Instead of making routing decisions for each packet **hop-by-hop**, UMW **dynamically chooses** routes of each packet **right** at the **source**.
- Instead of taking control actions based on queue lengths $Q(t)$ (**closed-loop control**), UMW is oblivious to the queues (**semi open-loop control**).

What's new? (Preview)

The new policy is named **Universal Max-Weight** Algorithm (UMW). It is different from the classical algorithms (such as, *Backpressure*) in the following aspects:

- Instead of making routing decisions for each packet **hop-by-hop**, UMW **dynamically chooses** routes of each packet **right** at the **source**.
- Instead of taking control actions based on queue lengths $Q(t)$ (**closed-loop control**), UMW is oblivious to the queues (**semi open-loop control**).
- UMW maintains a m -dimensional vector at the source, called **virtual queue-lengths** $\tilde{Q}(t)$, which are used for making **routing and scheduling decisions**.

What's new? (Preview)

The new policy is named **Universal Max-Weight** Algorithm (UMW). It is different from the classical algorithms (such as, *Backpressure*) in the following aspects:

- Instead of making routing decisions for each packet **hop-by-hop**, UMW **dynamically chooses** routes of each packet **right** at the **source**.
- Instead of taking control actions based on queue lengths $Q(t)$ (**closed-loop control**), UMW is oblivious to the queues (**semi open-loop control**).
- UMW maintains a m -dimensional vector at the source, called **virtual queue-lengths** $\tilde{Q}(t)$, which are used for making **routing and scheduling decisions**.
- The virtual queues $\tilde{Q}(t)$ correspond to a **precedence-relaxed** network with dynamics given by a **Lindley recursion**.


What's new? (Preview)

The new policy is named **Universal Max-Weight** Algorithm (UMW). It is different from the classical algorithms (such as, *Backpressure*) in the following aspects:

- Instead of making routing decisions for each packet **hop-by-hop**, UMW **dynamically chooses** routes of each packet **right** at the **source**.
- Instead of taking control actions based on queue lengths $Q(t)$ (**closed-loop control**), UMW is oblivious to the queues (**semi open-loop control**).
- UMW maintains a m -dimensional vector at the source, called **virtual queue-lengths** $\tilde{Q}(t)$, which are used for making **routing and scheduling decisions**.
- The virtual queues $\tilde{Q}(t)$ correspond to a **precedence-relaxed** network with dynamics given by a **Lindley recursion**.
- UMW uses the solution of some standard combinatorial problems (e.g., **Shortest Path, MST, Steiner Tree, MCDS..**) on a graph weighted by the dynamic virtual queues.

Design of UMW: Motivation and Insight


- **Observation:** Because of coupling, networked queues are **harder to analyze and control**. This is because queuing process, described by **Skorokhod maps** are fundamentally **non-linear** in nature.



IID arrivals to Q_1 causes **correlated arrivals** to Q_2

Design of UMW: Motivation and Insight

- **Observation:** Because of coupling, networked queues are **harder to analyze and control**. This is because queuing process, described by **Skorokhod maps** are fundamentally **non-linear** in nature.




IID arrivals to Q_1 causes **correlated arrivals** to Q_2

- This motivates us to obtain a **relaxed system**, which is easier to analyze, yet, preserves some fundamental characteristics we are interested in (e.g., **stability**) .

Design of UMW: Motivation and Insight

- **Observation:** Because of coupling, networked queues are **harder to analyze and control**. This is because queuing process, described by **Skorokhod maps** are fundamentally **non-linear** in nature.




IID arrivals to Q_1 causes **correlated arrivals** to Q_2

- This motivates us to obtain a **relaxed system**, which is easier to analyze, yet, preserves some fundamental characteristics we are interested in (e.g., **stability**) .

Question: How to obtain a **good** relaxation ? Which constraints to relax ?

Design of UMW: Motivation and Insight

- **Observation:** Because of coupling, networked queues are **harder to analyze and control**. This is because queuing process, described by **Skorokhod maps** are fundamentally **non-linear** in nature.




IID arrivals to Q_1 causes **correlated arrivals** to Q_2

- This motivates us to obtain a **relaxed system**, which is easier to analyze, yet, preserves some fundamental characteristics we are interested in (e.g., **stability**) .


Question: How to obtain a **good** relaxation ? Which constraints to relax ?

Ans: The Precedence Constraints!

Precedence Relaxation: Example




Precedence Relaxation: Example




$$\text{path}^* = \{\{1, 2\}, \{2, 3\}, \{3, 4\}\}$$

Precedence: The packet p **cannot** be transmitted over the link $2 - 3$, until it has been transmitted over the link $1 - 2$.

Precedence Relaxation: Example



$$\text{path}^* = \{\{1, 2\}, \{2, 3\}, \{3, 4\}\}$$



Virtual Queues

Virtual Net: Packets are replicated to the virtual queues as soon as they arrive to the source.

Virtual Queues: Operation

- Associate a virtual queue with each link of the graph.
- Upon packet arrival:
 - Determine an *optimal route* $T^*(t)$ (e.g., path, spanning tree) for the packet
 - *Immediately* inject a new virtual packet to each virtual queue along the route
 - This amounts to incrementing the queue counters along the route
- Serve the virtual packets at the rate $\mu^*(t)$ as long as the corresponding virtual queues are non-empty
 - Subject to the same link scheduling constraints ($\mu^*(t) \in \mathcal{M}$)

Virtual Queues: Operation

- Associate a virtual queue with each link of the graph.
- Upon packet arrival:
 - Determine an *optimal route* $T^*(t)$ (e.g., path, spanning tree) for the packet
 - *Immediately* inject a new virtual packet to each virtual queue along the route
 - This amounts to incrementing the queue counters along the route
- Serve the virtual packets at the rate $\mu^*(t)$ as long as the corresponding virtual queues are non-empty
 - Subject to the same link scheduling constraints ($\mu^*(t) \in \mathcal{M}$)

Question: How to design optimal controls: $T^*(t)$ and $\mu^*(t)$?

Dynamics of the Virtual Queues $\tilde{Q}(t)$

The virtual queues can be mathematically identified with an m -dimensional vector taking values in $\mathbb{Z}_+^{|E|}$.

Dynamics of the Virtual Queues $\tilde{Q}(t)$

The virtual queues can be mathematically identified with an m -dimensional vector taking values in $\mathbb{Z}_+^{|E|}$.

► Denote the (controlled) arrival to the VQ \tilde{Q}_e by $\tilde{A}_e(t)$. Then, the virtual queues follow the following Lindley dynamics:

$$\tilde{Q}_e(t+1) = (\tilde{Q}_e(t) + \tilde{A}_e(t) - \mu_e(t))_+, \quad (\text{Lindley recursion}) \quad (1)$$

► Note that, the arrivals to the virtual queues $(A_e(t), e \in E)$ are **explicit control variables** at the source.

► Unlike the original system, given the controls, the virtual queues are **independent** of each other. This makes exact analysis tractable.

Stabilizing Controls for $\tilde{Q}(t)$: Drift Analysis

- We design the per-slot controls $\pi^{\text{UMW}} \equiv (\mathbf{A}(t), \boldsymbol{\mu}(t))_{t \geq 0}$, **stabilizing** the virtual system $\{\tilde{Q}(t)\}_{t \geq 0}$.
- The policy consists of the routing decisions : **routing** $\mathbf{A}^\pi(t)$, sand **scheduling** $\boldsymbol{\mu}^\pi(t)$.
- Intuition: This control is *likely to stabilize* the physical queues as well
 - However, note that the dynamics of the physical queues depend explicitly on the packet scheduling policy (e.g., FIFO, LIFO etc.)

Stabilizing Controls for $\tilde{Q}(t)$: Drift Analysis

- We design the per-slot controls $\pi^{\text{UMW}} \equiv (\mathbf{A}(t), \boldsymbol{\mu}(t))_{t \geq 0}$, **stabilizing** the virtual system $\{\tilde{Q}(t)\}_{t \geq 0}$.
- The policy consists of the routing decisions : **routing** $\mathbf{A}^\pi(t)$, and **scheduling** $\boldsymbol{\mu}^\pi(t)$.
- Intuition: This control is *likely to stabilize* the physical queues as well
 - However, note that the dynamics of the physical queues depend explicitly on the packet scheduling policy (e.g., FIFO, LIFO etc.)
- To stabilize the virtual queues, we choose the control that **minimizes the drift of the Quadratic Lyapunov function of the Virtual Queues**.
- ► It turns out that the controls $\mathbf{A}(t)$ (routing) and $\boldsymbol{\mu}(t)$ (activation) are **separable** and are **standard combinatorial problems** that can be solved at the source r.

Derivation of the Control-Policy

- Define a Quadratic Lyapunov function

$$L(\tilde{\mathbf{Q}}(t)) \stackrel{\text{def}}{=} \sum_{e \in E} \tilde{Q}_e^2(t)$$

- The one-slot drift of $L(\tilde{\mathbf{Q}})(t)$ under any admissible policy π may be computed to be

$$\begin{aligned} \Delta^\pi(t) &\stackrel{\text{def}}{=} L(\tilde{\mathbf{Q}}(t+1)) - L(\tilde{\mathbf{Q}}(t)) \\ &\leq B + 2 \left(\underbrace{\sum_{e \in E} \tilde{Q}_e(t) A(t) \mathbb{1}(e \in \mathcal{T}^\pi(t))}_{(a)} - \underbrace{\sum_{e \in E} \tilde{Q}_e(t) \mu_e^\pi(t)}_{(b)} \right) \quad (2) \end{aligned}$$

Where $\mathcal{T}^\pi(t) \in \mathcal{T}$ and $\mu^\pi(t) \in \mathcal{M}$ are routing and activation control variables chosen for slot t .

- The drift upper-bound (2) has a nice separable form and may be minimized over the set of all feasible controls individually.

Optimal Routing Policy $T^*(t)$

Minimizing the term (a), we get the following optimal routing policy.

Optimal Routing : $T^*(t)$

$$T^*(t) \in \arg \min_{T \in \mathcal{T}} \sum_{e \in E} \tilde{Q}_e(t) \mathbb{1}(e \in T)$$

Optimal Routing Policy $T^*(t)$

Minimizing the term (a), we get the following optimal routing policy.


Optimal Routing : $T^*(t)$

$$T^*(t) \in \arg \min_{T \in \mathcal{T}} \sum_{e \in E} \tilde{Q}_e(t) \mathbb{1}(e \in T)$$

Examples:


- ▶ For the unicast problem : $T^*(t)$ is the **Shortest $s \rightarrow t$ path** in the weighted graph $\mathcal{G}(V, E, \tilde{Q}(t))$.
- ▶ For the broadcast problem : $T^*(t)$ is the **Minimum Weight Spanning tree (MST)** in the weighted graph $\mathcal{G}(V, E, \tilde{Q}(t))$.
- ▶ For the multicast problem : $T^*(t)$ is the **Minimum Weight Steiner tree** in the weighted graph $\mathcal{G}(V, E, \tilde{Q}(t))$ connecting the source nodes to the destination nodes.

Example of Optimal Routing: Unicast



Network Weighted by Virtual Queue sizes


Example of Optimal Routing: Unicast



Shortest 1-4 path = $\{\{1, 2\}, \{2, 4\}\}$

Network Weighted by Virtual Queue sizes

Example of Optimal Routing: Broadcast



MST rooted at 1 = $\{\{1, 2\}, \{2, 3\}, \{2, 4\}\}$

Network Weighted by Virtual Queue sizes

Optimal Link Activations $\mu^*(t)$


Similarly, minimizing the term (b), we obtain the following optimal activation policy

Optimal Activation: $\mu^*(t)$

$$\mu^*(t) \in \arg \min_{\mu \in \mathcal{M}} \sum_e \tilde{Q}_e(t) c_e \mathbb{1}(e \in \mu)$$


Example: For the case of wireless network with primary interference constraint, this problem corresponds to finding **Max-Weight-Matching** in the graph $\mathcal{G}(V, E, \mathbf{c} \odot \tilde{\mathbf{Q}}(t))$.

Example of Optimal Link Activations



Network Weighted by Virtual Queue sizes

Example of Optimal Link Activations



Max Weight Matching = $\{\{1, 3\}, \{2, 4\}\}$

Network Weighted by Virtual Queue sizes

Stability of the Virtual Queue

Theorem 6: Strong Stability of $\tilde{Q}(t)$

Under the above routing and scheduling policy, for all arrival rate $\lambda \in \mathring{\Lambda}$ the virtual queue process is Strongly stable and has a limiting M.G.F, i.e.,

$$\limsup_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \sum_e \mathbb{E}(\tilde{Q}_e(t)) \leq B$$

and,

$$\limsup_{T \rightarrow \infty} \mathbb{E}(\exp(\theta^* \sum_e \tilde{Q}_e(t))) \leq C$$

for some finite B, C and strictly positive θ^* .

Stability of the Virtual Queue

Theorem 6: Strong Stability of $\tilde{Q}(t)$

Under the above routing and scheduling policy, for all arrival rate $\lambda \in \tilde{\Lambda}$ the virtual queue process is Strongly stable and has a limiting M.G.F, i.e.,

$$\limsup_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \sum_e \mathbb{E}(\tilde{Q}_e(t)) \leq B$$

and,

$$\limsup_{T \rightarrow \infty} \mathbb{E}(\exp(\theta^* \sum_e \tilde{Q}_e(t))) \leq C$$

for some finite B, C and strictly positive θ^* .

The above leads to the following :

Lemma: Sample Path bound on Virtual Queues

Under the same condition, we have

$$\sum_e \tilde{Q}_e(t) = \mathcal{O}(\log t), \quad \text{a.s.}$$

Relating Virtual Queues to Physical Processes : Arrivals and Service

For any edge $e \in E$, define the **cumulative arrival and service processes** in the interval $(t_1, t_2]$ as follows

$$\tilde{A}_e([t_1, t_2]) \stackrel{(\text{def})}{=} \sum_{\tau=t_1}^{t_2} A_e(\tau), \quad \tilde{S}_e([t_1, t_2]) \stackrel{(\text{def})}{=} \sum_{\tau=t-1}^{t_2} \mu_e(\tau)$$

Relating Virtual Queues to Physical Processes : Arrivals and Service

For any edge $e \in E$, define the **cumulative arrival and service processes** in the interval $(t_1, t_2]$ as follows

$$\tilde{A}_e([t_1, t_2]) \stackrel{(\text{def})}{=} \sum_{\tau=t_1}^{t_2} A_e(\tau), \quad \tilde{S}_e([t_1, t_2]) \stackrel{(\text{def})}{=} \sum_{\tau=t-1}^{t_2} \mu_e(\tau)$$

Then the Lindley recursion takes the following form (**Skorokhod Map**):

$$\tilde{Q}_e(t) = \sup_{0 \leq \tau \leq t} (\tilde{A}_e(\tau, t) - \tilde{S}_e(\tau, t))$$

Thus, we have for any $0 \leq \tau \leq t$:

$$\tilde{A}_e(\tau, t) \leq \tilde{S}_e(\tau, t) + F(t), \quad F(t) = \mathcal{O}(\log t) \tag{3}$$

Optimal Control : Packet Scheduling

- How do we decide which packet to transmit over a link at any given time slot?
 - Why does it matter? Cant we just use FCFS?
- Nearest to Origin (NTO) policy [Gamarnik, 1998]
- **Extended Nearest to Origin policy (ENTO)**: When multiple packets contend for an edge, schedule the one which has traversed the least number of edges
 - Extension of NTO to general flow problems

Optimal Control : Packet Scheduling


- How do we decide which packet to transmit over a link at any given time slot?
 - Why does it matter? Cant we just use FCFS?
- Nearest to Origin (NTO) policy [Gamarnik, 1998]
- **Extended Nearest to Origin policy (ENTO)**: When multiple packets contend for an edge, schedule the one which has traversed the least number of edges
 - Extension of NTO to general flow problems

Theorem 7: Stability of the Physical Queues

The overall UMW policy is throughput-optimal.

Proof uses the previous **almost sure** arrival bound on a typical sample path with an **inductive argument** on the edges.

ENTO: Example





Packet p_1 has higher priority than p_2 to cross e_3 as it has traversed **less** number of edges



Outline

- 1 Introduction
- 2 Optimal Broadcasting in a DAG
- 3 Algorithm for Generalized Flow
- 4 Simulation Results**
- 5 Conclusion


Multi-commodity Flow Simulation



Broadcasting Simulation: Static Network



Broadcasting: Time Varying Network



Comparison of the time-averaged total queue lengths under the optimal (solid line) and heuristic (dashed line) UMW policy in the time-varying grid network (with parameter p_{ON}), for the broadcast problem.

Generality of the Virtual Queue Paradigm: Most Recent Results

- The algorithmic paradigm that we introduced in our thesis turns out to be pretty [general](#).
- Recently, we developed an optimal broadcast policy for wireless networks with [point-to-multi-point](#) links
- Here instead of trees, the “routes” consists of [Minimum Connected Dominating Sets](#) (MCDS)
- This framework could also be extended to [Network Utility Maximization](#) Problems.

Outline

- 1 Introduction
- 2 Optimal Broadcasting in a DAG
- 3 Algorithm for Generalized Flow
- 4 Simulation Results
- 5 Conclusion**

Conclusion

- Our understanding of network control theory has progressed enormously over the past 25 years, starting with the seminal work of Tassiulas and Ephremides in 1992.
- **Universal-Max-Weight** (UMW) is throughput-optimal and can be used in a wide range of network flow problems
- UMW solves the important open problem of **Generalized Flow** and eliminates many of the drawbacks of the classical Back-Pressure algorithm
- **Open Problem:** Does the algorithm remain optimal when used in conjunction with the Physical Queue lengths, instead of the virtual queue-lengths ?
 - Empirical evidence suggest **yes**
 - Leads to a more efficient practical implementation of UMW

References

Accepted:

- ① **A. Sinha** and E. Modiano, Optimal control for generalized network-flow problems, in IEEE [INFOCOM '17](#).
- ② **A. Sinha**, G. Paschos, C. P. Li, and E. Modiano, Throughput-optimal multihop broadcast on directed acyclic wireless networks, in [IEEE/ACM Transactions on Networking](#)
- ③ **A. Sinha**, L. Tassiulas, and E. Modiano, Throughput-optimal broadcast in wireless networks with dynamic topology, in ACM [MobiHoc '16](#) (**Best Paper Award**).
- ④ **A Sinha**, G. Paschos, and E. Modiano, Throughput-optimal multi-hop broadcast algorithms, in ACM [MobiHoc '16](#)
- ⑤ **A. Sinha**, G. Paschos, C.-p. Li, and E. Modiano, Throughput-optimal broadcast on directed acyclic graphs, in IEEE [INFOCOM '15](#)

Under submission:

- ① **A. Sinha** and E. Modiano, Optimal control for generalized network-flow problems, in [IEEE/ACM Transactions on Networking](#)
- ② **A. Sinha** and E. Modiano, Throughput-Optimal Broadcast in Wireless Networks with Point-to-Multipoint Transmissions, in ACM [MobiHoc '17](#),
- ③ **A. Sinha**, L. Tassiulas, and E. Modiano, Throughput-optimal broadcast in wireless networks with dynamic topology, in [IEEE Trans. on Mobile Computing](#)
- ④ **A Sinha**, G. Paschos, and E. Modiano, Throughput-optimal multi-hop broadcast algorithms, in [IEEE/ACM Transactions on Networking](#)